

Optimize and speed up U-Net for Generalized Nuclear Segmentation for Computational Pathology

Bui Kim Dung, Quan Yu, Tan Xinli Steven, Meng Yang and Lwi Tiong Chai

Abstract—Nuclear Segmentation is the extraction of the cell boundary locations of the cells in the images taken of the tissue samples, a key step in extracting biomedical data from tissue samples. There are many challenges to extracting the cell boundary locations from the images due to the complex, crowded images and similarities in the appearance between each cell. This paper used the data set from the MoNuSeg Nuclear Segmentation challenge, which was taken from multiple hospitals and includes a diversity of nuclear appearances from several patients, disease states, and organs, techniques trained on it are likely to generalize well and work right out-of-the-box on other H&E-stained images. This paper summarizes the related work related to deep learning-based nuclear segmentation, then based on the state-of-the-art models, it proposes an optimized architecture of U-net Convolutional Network calls All Convolutional U-Net to optimize and increase nuclei segmentation result while speeding up the nuclear segmentation model training process.

Index Terms—Annotation, boundaries, dataset, deep learning, U-Net, ACU-Net, All convolutional u-net, convolutional neural network, nuclear segmentation, nuclei.

I. INTRODUCTION

THIS paper aims to improve the process of nuclear segmentation of the cells for the MoNuSeg Nuclear Segmentation Challenge. The MoNuSeg is an official satellite event of MICCAR 2018, the 21st International Conference on Medical Image Computing and Computer Assisted Intervention. Nuclear segmentation is the process of identifying the boundaries of the cells in the digital image accurately and precisely by partitioning a digital image of the cells into multiple segments.

Cell samples are taken from patients and stained with Haematoxylin and eosin (H&E). Haematoxylin renders nuclear dark bluish purple and epithelium light purple, while eosin renders stroma pink. This makes it easier for people and computers to identify the cell boundaries and nuclei in the cells. Studying the size, density, shape and other features of the cells, medical conditions such as cancer or sclerosis can be diagnosed and appropriate treatments can be recommended. Furthermore,

studying the shape and size of the cells and nuclei can help give a better understanding of the human body. Nuclear Segmentation can also be applied to cell samples taken from plants and animals. In summary, the quantitative assessment of the size and shape of the cells has many significant research and industrial applications in the biomedical field.

II. DATA SETS

A. Background

The training images were taken from the MoNuSeg Nuclear Segmentation challenge 2018, (<https://monuseg.grand-challenge.org>).

The dataset consists of tissue samples taken from several patients with tumours of different organs. The tissue samples were stained with H&E and captured at 40x magnification. As a result, these tissue samples have very different shapes and sizes, which should result in a more generalised nuclear segmentation algorithm.

The dataset also includes the manually annotated nuclear boundaries of the cells in the images, which can be used to assess the accuracy of the Nuclear Segmentation algorithm. These were included as xml files.

B. Data Preprocessing

1) Pre-process annotation files

The binary mask images of the nuclear boundaries were pre-processed and two-class mask images with the pixels inside the nuclear boundaries were filled using the OpenCV-Python algorithm. First, morphological dilation and closing were applied using `cv2.dilate` and `cv2.MORPH_CLOSE` functions consecutively. After the transformation of the nuclear boundaries into contours without gaps, function `binary_fill_holes` from Scipy was used to fill the pixels inside the nuclear boundaries, resulting in usable mask images with pixels inside the nuclei labeled. The result of dilation + closing and binary filling holes is shown in Fig 1.

2) Create training patches

The original dataset consisted of 30 images of 1000 x 1000




Fig. 1. Two-class mask image generation

resolution. To increase the training data and reduce the computational complexity, the images were split into 64x64 pixels images using a sliding window with stride of 30 pixels to create the new training sets with total 24025 sample tiles. The 64x64 image size corresponds to the size of the largest nuclei, while the stride of 30 pixels create overlapped edges among the samples to prevent data loss as well as to produce a training set with higher resolution.

III. RELATED WORK: DEEP LEARNING-BASED NUCLEAR SEGMENTATION

There are countless varieties of Convolutional Neural Networks (CNN) that have been used for image classification, and many of them have high accuracy and precision. [1]. For bio-medical images segmentation, the most popular technique is to apply a CNN on high resolution images to produce probability maps for detecting various tissue segments based on processing small patches of fixed sizes that have been sampled from the large original images.

A. Two-class CNN

The two-class CNN model tries to classify whether the central pixel of a patch belongs to a nucleus or not [2]. The result when applying this model to all possible overlapping patches of a fixed size is a binary or nuclear probability map. To resolve touching nuclei, it has been proposed that distance transform be computed on the binary map, which is likely to yield two separate local minimas near the centers of two touching nuclei due to the concavity in the shape of multiple nuclei. M. E. Plissiti, C. Nikou [3] and M. Veta, et al [4] in their experiments on automatic nuclei segmentation have also used shape concavity to separate touching nuclei. This technique works well for simple cases of two touching nuclei but not for more complex cases [5].

B. Three-Class CNN Emphasizing Nuclear Boundaries

Neeraj Kumar et al. [5] proposed the Three-Class CNN (CNN3). By explicitly introducing the nuclear *boundary* as a third class of pixels in addition to the usual binary of foreground (*inside* any nucleus) and background (*outside* every nucleus), this CNN produces a ternary probability map, which emphasizes detection of nuclear boundary and solves the

complex cases of touching nuclei. The architecture of the CNN3 is shown in Table I.

The strategy in Neeraj Kumar et al. [5] has two drawbacks. First, it is quite slow because to estimate the 3-class probability assignment on each pixel, each 3-output node CNN took a patch of size 51 x 51 centered at that pixel as input. The network must be run separately on each patch (total 171,000 patches), and there is a lot of redundancy due to overlapping patches.

Secondly, the trade-off between localization accuracy and

TABLE I
CNN3 ARCHITECTURE FOR SEGMENTING NUCLEI

Layer	Filter size	Activation	Output size	Dropout rate
Input	-		51 x 51 x 3	
Conv 1	4 x 4	ReLU	48 x 48 x 25	0.1
Pool 1	2 x 2	Max	24 x 24 x 25	-
Conv 2	5 x 5	ReLU	20 x 20 x 50	0.2
Pool 2	2 x 2	Max	10 x 10 x 50	-
Conv 3	6 x 6	ReLU	5 x 5 x 80	0.25
Pool 3	2 x 2	Max	3 x 3 x 80	-
FC 1	-	ReLU	1024	0.5
FC2	-	ReLU	1024	0.5
Output	-	SoftMax	3	

the use of context prevents the model from achieving better results. The size 51 x 51 was chosen as the smallest size to cover most of the large nuclei and to reduce computation complexity and training time. Patch-wise training is common [6] but lacks the efficiency of fully convolutional training.

This architecture gave 92% accuracy for binary pixel classification and 0.76 average Dice's coefficient, compared to [8] which gave only 78% accuracy, CNN3 set the new benchmarking for generalized nuclear segmentation techniques.

C. Fully Convolutional Networks

Adapted contemporary classification networks (AlexNet




Fig. 2. FCN skip architecture which learns to combine coarse, high layer information with fine, low layer information. Pooling and prediction layers are shown as grids that reveal relative spatial coarseness, while intermediate layers are shown as vertical lines. [6]

[18], the VGG net [19], and GoogLeNet [20]) into *fully convolutional networks (FCN)* [6], a group of authors from UC Berkeley has achieved state-of-the-art segmentation of PASCAL VOC. The model has been the first work to train FCNs end-to-end for pixelwise prediction and from supervised pre-training.

By fine-tuning to the segmentation task, the model transfers the learned representations from the classical classification networks to dense prediction.

Semantic segmentation faces an inherent tension between semantics and locations: global information resolves what while local information resolves where [6]. Deep features hierarchies encode location and semantics in a nonlinear local-to-global pyramid. The skip architecture introduced in this paper as shown in Fig 2. that combines semantic information from a deep, coarse layer with appearance information from a shallow, fine layer to produce accurate and detailed segmentations.

However, although proven to achieve state-of-the-art segmentation accuracy, the deep architecture of these fully convolutional networks requires large amount of training input to produce good accuracy and generalize well on unseen samples, while most of the bio-medical segmentation problems have very limited number of sample images [5].

D. U-Net Convolutional Network

In 2015, Olaf Ronneberger, Philipp Fischer, and Thomas Brox introduced U-Net Convolutional Network [7] on the ISBI




Fig. 3. U-net architecture. Each blue box corresponds to a multi-channel feature map. The number of channels is denoted on top of the box. The x-y-size is provided at the lower left edge of the box. White boxes represent copied feature maps. The arrows denote the different operations. [7]

challenge for segmentation of neuronal structures in electron microscopic stacks.

U-net has been built upon the fully convolutional network [6] architecture, with modification to adapt it to very few training images, while still yield precise segmentations. The architecture of U-net is depicted in Fig 3.

The important modification in U-net architecture is that in the up-sampling part, it still maintains large number of feature channels, which allows the network to propagate context information to higher resolution layers. The expansive path is symmetric to the contracting path and yields a u-shaped architecture.

The original U-net model produces total more than 31 million parameters, using the input images of size 512x512 pixels, training of 10 hours on a Nvidia Titan GPU (6GB).

IV. PROPOSED MODELS: CUSTOMIZED U-NET FOR GENERALIZED NUCLEAR SEGMENTATION

This paper introduces two customized architectures of U-net which speed up the training process and is optimized for the generalized nuclei segmentation task on very small datasets. Both models were able to achieve higher average Dice's coefficient compared to the state-of-the-art CNN3 after only fewer training epochs which took 41 minutes (compared to the 16.5 hours for training CNN3) on Nvidia GTX 1070 (8GB).

A. Network Architecture

The overview architecture of one of the proposed architectures is depicted in Fig 4.

Similar to the original U-net, it consists of a contracting path (left side) and an expansive path (right side). The contracting path follows the typical architecture of convolutional network, which consists of the repeated application of two 3x3 convolutions. At each down-sampling step, the number of feature channels is doubled.

1) All Convolutional U-Net

The first customized model is a homogeneous network which solely consists of convolutional layer. The architecture is depicted in Fig 4. It is different from the original U-net in 3 major configurations:

First, in the down-sampling path, all the max pooling layers are replaced by 2x2 convolutional layers with stride 2x2 with ReLU activation. The idea of replacing max pooling layer by a normal convolutional layer was first introduced in [22] and proved to be more efficient without loss in accuracy rate. This replacement increases the number of trainable parameters, force the network to learn to sub-sample the input instead of simply pooling. Experiments to compare the same network with and without the max pooling replacement on the same dataset showed an improvement of Dice coefficient score by 0.03.

Secondly, except the sub-sampling convolutional layers which replaces max-pooling, all the convolutional layers are using same padding, while the original U-Net does not use padding. This different is because while the U-net use a much bigger resolution input (572x572), the proposed architecture prefers smaller input size to reduce computation complexity (64x64 is chosen since this is the smallest size which can cover most of large nuclei). Another reason for the same padding is because during the preprocessing step, the original 1000x1000 images are cropped into multiple 64x64, same padding convolutional will help to reserve the border region of the small tiles and reserve the context due to cropping.

Thirdly, in the up-sampling path, while original U-net uses up-sampling 2D layer to extrapolate the border, then crops the image from left side (the skip connection) to sync the input shape, then concatenate with the up-sampling output. This model uses transpose up-convolution layer to double the size of the input, then simply concatenate with the output from left side




Fig. 4. Proposed architecture. Each blue box corresponds to a multi-channel feature map. The number of channels is denoted on top of the box. The x-y-size is provided at the lower left edge of the box. White boxes represent copied feature maps. The arrows denote the different operations. The setting of channel size is the setting of the lightweight version. The details settings are described at Table II.

(skip connection arrows). This modification also contributes to the reduce of training time and computation complexity.

2) Customized U-net with ELU

The team took another approach to optimize U-net for the nuclei segmentation problem. While still keeping the max pooling and up-sampling layers as in original U-net, the second customized model make use of the exponential linear units in all the convolutional layers.

Similar to the proposed All Convolutional U-net at IV.A.1, all the convolutional layers use same padding to reserve the border region of the small tiles and reserve the context due to cropping.

Table II shows the details configuration of this customized model.

B. Metrics

1) Loss function

This architecture does not support the three-class pixels as in CNN3. To address the challenge of separation of nuclei touching borders, the model makes use of weighted loss, where the separating background labels between touching cells is used to obtain a large weight in the loss function.

The binary cross entropy loss is calculated as:

$$-(y \log(p) + (1 - y) \log(1 - p)) \quad (1)$$

In which, y is the binary indicator (0 or 1) and p is the predicted probability observation o is of class C .

2) Evaluation Metric

A commonly used object detection metric is the F1-score. For ground true objects G_i indexed by i and segmented objects S_j indexed by j , the F1-score is based on true positives TP (count of all ground truth objects G_i with an associated segmented (detected) object S_j , false positive FP (count of all segmented objects S_j without a corresponding ground truth object G_i), and false negatives FN (count of all ground truth objects G_i without a corresponding detected object S_j). F1-score is defined as follows:

$$F1 = \frac{2TP}{2TP + FP + FN} \quad (2)$$

A major shortcoming of F1-score using any object association criteria is that it does not take pixel-level (segmentation) errors into account. Thus, for this model, beside accuracy metric, it used Average Dice's coefficient metric to report shape concordance (pixel-level errors) between the ground truth objects and their associated segmented objects.

$$Dice = \frac{2|G_i \cap S_j|}{|G_i| + |S_j|} \quad (3)$$

Another reason the Dice's coefficient metric was chosen was to make it easier to compare with the result from the baseline model CNN3 [5].

C. Training and testing

Both proposed networks can be trained on end-to-end by backward propagation and stochastic gradient descent (SGD). The total sample set was split into training set: evaluation set: test set with ratio 0.7:0.1:0.2 with random shuffling, using Adam optimizer and batch size of 32 samples.

Nuclei were detected by thresholding the class probability map at 0.5.

One big advantage observed here is that all the proposed models could achieved higher score without the need of complicated data augmentation on the training set, as well as any complex post processing techniques.

The training process also took significantly less time (less than 1 hour) compare to 16 hours training for the CNN3.

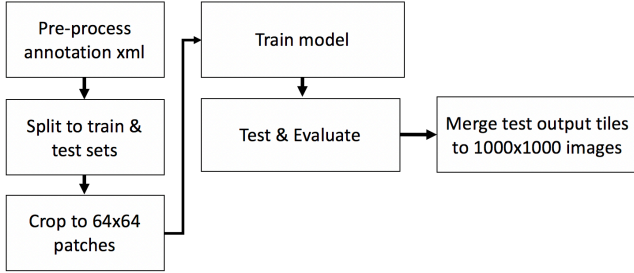


Fig. 5. Training and testing process

V. EXPERIMENTAL RESULTS

A. All Convolutional U-net

The input size for this model is 64x64x3, trained with batch size of 32 samples over total 24025 samples. With total 32,425,345 parameters, this model took 30 minutes over 18 epochs on Nvidia GTX 1070 (8GB).

It was able to achieve 0.8245 average Dice's coefficient (while the CNN3 achieved 0.76).

The sample test images are shown at Fig 6.

B. Customized U-net with ELU

Details configuration of this model are described in table II. The input size for this model is 64x64x3, trained with batch size of 32 samples over total 24025 samples.

This model achieves Dice's coefficient of 0.8349 and 91.4% accuracy. Total time taken to train this model is 41 minutes on Nvidia GTX 1070 (8GB) with 25 epochs.

C. Fine-tuning

To investigate further performance, the team added a callback function after every epoch and found out that both models achieved high training Dice's coefficient at very early stages (after only the 4th training epoch in case of the ELU U-net). Both models also seem to be overfit at some of the first round of training, the team had to add Drop out in between of all the 2 consecutive convolutional layers in both down-sampling and up-sampling paths.

The Keras-based (TensorFlow backend) [21][24]

TABLE II
ARCHITECTURE OF U-NET WITH EXPONENTIAL LINEAR UNITS

Layer	Filter size	Activation	Output size	Padding
Input	-	-	64 x 64 x 3	-
Conv 1	3 x 3	ELU	64 x 64 x 64	same
Conv 2	3 x 3	ELU	64 x 64 x 64	same
Pool 1	2 x 2	-	32 x 32 x 64	-
Conv 3	3 x 3	ELU	32 x 32 x 128	same
Conv 4	3 x 3	ELU	32 x 32 x 128	same
Pool 2	2 x 2	-	16 x 16 x 128	-
Conv 5	3 x 3	ELU	16 x 16 x 256	same
Conv 6	3 x 3	ELU	16 x 16 x 256	same
Pool 3	2 x 2	-	8 x 8 x 256	-
Conv 7	3 x 3	ELU	8 x 8 x 512	same
Conv 7	3 x 3	ELU	8 x 8 x 512	same
Pool 4	2 x 2	-	4 x 4 x 512	-
Conv 8	3 x 3	ELU	4 x 4 x 1024	same
Conv 9	3 x 3	ELU	4 x 4 x 1024	same
Upsample+Concat	2 x 2	-	8 x 8 x 1024	-
Conv 10	3 x 3	ELU	8 x 8 x 512	same
Conv 11	3 x 3	ELU	8 x 8 x 512	same
Upsample+Concat	2 x 2	-	16 x 16 x 512	-
Conv 12	3 x 3	ELU	16 x 16 x 256	same
Conv 13	3 x 3	ELU	16 x 16 x 256	same
Upsample+Concat	2 x 2	-	32 x 32 x 256	-
Conv 14	3 x 3	ELU	32 x 32 x 128	same
Conv 15	3 x 3	ELU	32 x 32 x 128	same
Upsample+Concat	2 x 2	-	64 x 64 x 128	-
Conv 16	3 x 3	ELU	64 x 64 x 64	same
Conv 17	3 x 3	ELU	64 x 64 x 64	same
Output	1 x 1	Sigmoid	64 x 64 x 1	-

TABLE III
PERFORMANCE COMPARISON

Model	Dice's Coefficient	Epoch	Training time
CNN3	0.7623	115	16.5 hours (6GB GPU)
NB-model	0.812	300	7.5 hours
U-net Trans	0.7785	55	1.5 hours
All Conv U-net	0.8245	18	30 mins (8GB GPU)
U-net ELU	0.8349	25	41 mins (8GB GPU)

experimental implementation of both proposed models are available at https://github.com/kdung/cell_segmentation.

VI. CONCLUSION

Table III shows details of comparison among the 4 models: the baseline CNN3, the Nuclei-Boundary model [23], and the proposed All Convolutional U-net, the customized U-net with ELU. The U-net with ELU achieves the best performance. The training time comparison is only for reference since the same machine configuration as the baseline was not available.

By exploring the different techniques and trying to overcome the drawbacks of existing architecture in biomedical image segmentation, this paper proposes a customized architecture of U-net model, which proven to be able to achieve better segmentation results in significant less epochs, computation complexity, complex pre and post processing, and reduce




Fig. 6. Sample test results from the proposed model U-net with ELU activation. The pictures on the left most column are original cropped images from the dataset. The images in middle column (back and white images) are corresponding ground truth annotated nuclei. The right most column contains the predicted probability map segmented.

training time from hours to minutes compare to the baseline CNN3 [5]

ACKNOWLEDGMENT

The authors thank Dr. Matthew Chua and Dr. Tian Jing at National University of Singapore, Institute of Systems Science for their guidance and suggesting new experiments to strengthen the paper.

REFERENCES

- [1] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, pp. 436–444, May 2015.
- [2] Ian Goodfellow, Yoshua Bengio, and Aaron Courville, "Convolutional Networks," in *Deep Learning*, Cambridge: MIT Press, 2016, ch. 9, sec. 2, pp. 321–363.
- [3] M. E. Plissiti and C. Nikou, "Overlapping cell nuclei segmentation using a spatially adaptive active physical model," *IEEE Trans. Image Process.*, vol. 21, no. 11, pp. 4568–4580, Nov. 2012.
- [4] M. N. Gurcan, L. E. Boucheron, A. Can, A. Madabhushi, N. M. Rajpoot, and B. Yener, "Histopathological image analysis: A review," *IEEE Rev. Biomed. Eng.*, vol. 2, pp. 147–171, Oct. 2009.
- [5] N. Kumar, R. Verma, S. Sharma, S. Bhargava, A. Vahadane and A. Sethi, "Dataset and a Technique for Generalized Nuclear Segmentation for Computational Pathology," *IEEE Trans. Med. Imag.*, vol. 36, no. 7, Jul. 2017.
- [6] J. Long, E. Shelhamer and T. Darrell "Fully Convolutional Networks for Semantic Segmentation" UC Berkeley. Available: https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf [Accessed: 14 Aug 2018]
- [7] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation", 2015. Computer Science Department and BIOS Centre for Biological Signalling Studies, University of Freiburg, Germany. Available: <https://arxiv.org/pdf/1505.04597.pdf>. [Accessed on 14 Aug 2018]
- [8] F. Xing, Y. Xie, and L. Yang, "An automatic learning-based framework for robust nucleus segmentation," *IEEE Trans. Med. Imag.*, vol. 35, no. 2, pp. 550–566, Feb. 2016.
- [9] H. Llewellyn, "Observer variation, dysplasia grading, and HPV typing: A review," *Amer. J. Clin. Pathol.*, vol. 114, pp. S21–S35, Nov. 2000.
- [10] M. Veta, P. J. van Diest, R. Kornegoor, A. Huisman, M. A. Viergever, and J. P. W. Pluim, "Automatic nuclei segmentation in H&E stained breast cancer histopathology images," *PLoS ONE*, vol. 8, no. 7, p. e70221, 2013.
- [11] A. Vahadane and A. Sethi, "Towards generalized nuclear segmentation in histological images," in *Proc. IEEE 13th Int. Conf. Bioinform. Bioeng. (BIBE)*, Nov. 2013, pp. 1–4.

- [12] J. Odstrcilik et al., "Retinal vessel segmentation by improved matched filtering: Evaluation on a new high-resolution fundus image database," *IET Image Process.*, vol. 7, no. 4, pp. 373–383, Jun. 2013.
- [13] A. C. Ruifrok and D. A. Johnston, "Quantification of histochemical staining by color deconvolution," *Anal. Quant. Cytol. Histol.*, vol. 23, no. 4, pp. 291–299, Aug. 2001.
- [14] R. Collobert, S. Bengio, and J. Mariéthoz, "Torch: A modular machine learning software library," *Idiap Res. Inst., Martigny, Switzerland, Tech. Rep. EPFL-REPORT-82802*, 2002.
- [15] A. E. Carpenter et al., "CellProfiler: Image analysis software for identifying and quantifying cell phenotypes," *Genome Biol.*, vol. 7, no. 10, p. R100, 2006.
- [16] F. Dong et al., "Computational pathology to discriminate benign from malignant intraductal proliferations of the breast," *PLoS ONE*, vol. 9, no. 12, p. e114885, Dec. 2014.
- [17] J. Schindelin et al., "Fiji: An open-source platform for biological-image analysis," *Nature Methods*, vol. 9, no. 7, pp. 676–682, Jul. 2012.
- [18] A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. In *NIPS*, 2012
- [19] A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. In *NIPS*, 2012
- [20] C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich. Going deeper with convolutions. *CoRR*, abs/1409.4842, 2014.
- [21] Chollet, Francois and others, Keras, 2015, <https://keras.io>
- [22] Jost Tobias Springenberg*, Alexey Dosovitskiy*, Thomas Brox, Martin Riedmiller, "Striving for simplicity: The All Convolutional Net", Apr. 2015
- [23] Yuxin Cui*, Guiying Zhang*, Zhonghao Liu, Zheng Xiong, Jianjun Hu, Member, IEEE, "A Deep Learning Algorithm for One-step Contour Aware Nuclei Segmentation of Histopathological Images", Mar. 2018
- [24] Martin Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Rafal Jozefowicz, Yangqing Jia, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dan Mané, Mike Schuster, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org